


# Multi-agent Reinforcement learning with Logic Programming

Jie Meng


# Motivation

- We can't explain what AI agent have learned in a game
  - AlphaGo vs 李世石
- How to learn interpretable social rules
  - Combine RL and Logic rules



# Application Scenario

- Tic-Tac-Toe
- Escalator
- Traffic Rules



# Preliminary

- Logic Programming
- Logic Programming in Reinforcement Learning

# Preliminary: Logic Programming

- Inductive Logic Programming(ILP)

- ▶ clause:  $\alpha \leftarrow \alpha_1, \dots, \alpha_m$
- ▶ atom:  $p(t_1, t_2, \dots, t_n)$   
where  $p$  is a  $n$ -ary predicate and  $t_1, t_2, \dots, t_n$  are terms
- ▶ ground atom: an atom contains no variables

- ILP problem

- An ILP problem is a tuple  $(B; P; N)$  of ground atoms, where:

- ▶  $B$  is a set of background assumptions, a set of ground atoms.
    - ▶  $P$  is a set of positive instances - examples taken from the extension of the target predicate to be learned
    - ▶  $N$  is a set of negative instances - examples taken outside the extension of the target predicate

# Preliminary: Logic Programming

## An ILP problem

$$\mathcal{B} = \{zero(0), succ(0; 1), succ(1; 2), succ(2; 3), \dots\}$$
$$\mathcal{N} = \{even(1), even(3), even(5), even(7), \dots\}$$
$$\mathcal{P} = \{even(0), even(2), even(4), even(6), \dots\}$$

The aim of ILP is to construct a logic program that explains the positive instances and rejects the negative instances.

## One Solution

$$even(X) \leftarrow zero(X)$$
$$even(X) \leftarrow even(Y), succ2(Y, X)$$
$$succ2(X, Y) \leftarrow succ(X, Z), succ(Z, Y)$$

# Preliminary: Logic Programming in Reinforcement Learning

- Differentiable ILP

## Generating Atoms

Generating atoms from templates:

$$L = (\text{target Predicate}; \text{arity}; C)$$

## Valuations

Given a set  $G$  of  $n$  ground atoms, a valuation is a vector  $[0, 1]^n$  mapping each ground atom  $\gamma_i \in G$  to the real unit interval.

## Example

Language template:  $P_i = \{p/0, q/1, r/2\}; C = \{a, b\}$

$$\begin{array}{llll} \perp \mapsto 0.0 & p() \mapsto 0.0 & q(a) \mapsto 0.1 & q(b) \mapsto 0.3 \\ r(a, a) \mapsto 0.7 & r(a, b) \mapsto 0.1 & r(b, a) \mapsto 0.4 & r(b, b) \mapsto 0.2 \end{array}$$

# Preliminary: Logic Programming in Reinforcement Learning

## Generating Clauses

Generating atoms from templates:  $C_p^1, C_p^2, \dots, C_p^n$ , where  $C_p^i$  is a clause for target predicate P. e.g.  $p(X) \leftarrow r(Y), q(X, Y)$

## Weights

$W_p[j, k]$  represents how strongly the system believes that the pair of clauses  $(C_p^j, C_p^k)$  is the right way to define the intensional predicate p.

## Updating Valuations

$V^0 = (v_0, v_1, \dots, v_n) \rightarrow f_{infer}(clauses, weight, T)$   
→ conclusion valuations



# Preliminary: Logic Programming in Reinforcement Learning

## Frame

$$V^0 = \begin{bmatrix} b_1 \\ \dots \\ b_m \\ f_1 \\ \dots \\ f_n \end{bmatrix} = \begin{bmatrix} 1 \\ \dots \\ 1 \\ \dots \\ 0 \\ \dots \\ 0 \end{bmatrix} \Rightarrow f_{infer} \left( \begin{bmatrix} c_1 & c_2 \\ \dots & \dots \\ c_k & c_m \end{bmatrix}, \mathbf{W}, T \right) \Rightarrow V^T \rightarrow \begin{pmatrix} prob(n_i) \\ \dots \\ prob(p_i) \end{pmatrix}$$

## Loss

Calculating the probability of the label  $\lambda$  given the atom  $\alpha$ .

$$p(\lambda|\alpha, \mathbf{W}, \Pi, \mathcal{L}, \mathcal{B}) = f_{\text{extract}} \left( f_{\text{infer}} \left( f_{\text{convert}}(\mathcal{B}), f_{\text{generate}}(\Pi, \mathcal{L}), \mathbf{W}, T \right), \alpha \right)$$

$$loss = - \mathbb{E}_{(\alpha, \lambda) \sim \Lambda} [\lambda \cdot \log p(\lambda|\alpha, \mathbf{W}, \Pi, \mathcal{L}, \mathcal{B}) + (1 - \lambda) \cdot \log(1 - p(\lambda|\alpha, \mathbf{W}, \Pi, \mathcal{L}, \mathcal{B}))]$$

# Preliminary: Logic Programming in Reinforcement Learning

## MDP with logic interpretation

A triple  $(M, p_S, p_A)$ :

- ▶  $p_A : [0, 1]^{|D|} \rightarrow [0, 1]^{|A|}$ . Maps the valuation (or score) of a set of atoms  $D$  to the probability of actions.
- ▶  $p_S : S \rightarrow 2^G$ . Maps each states to ground atoms.

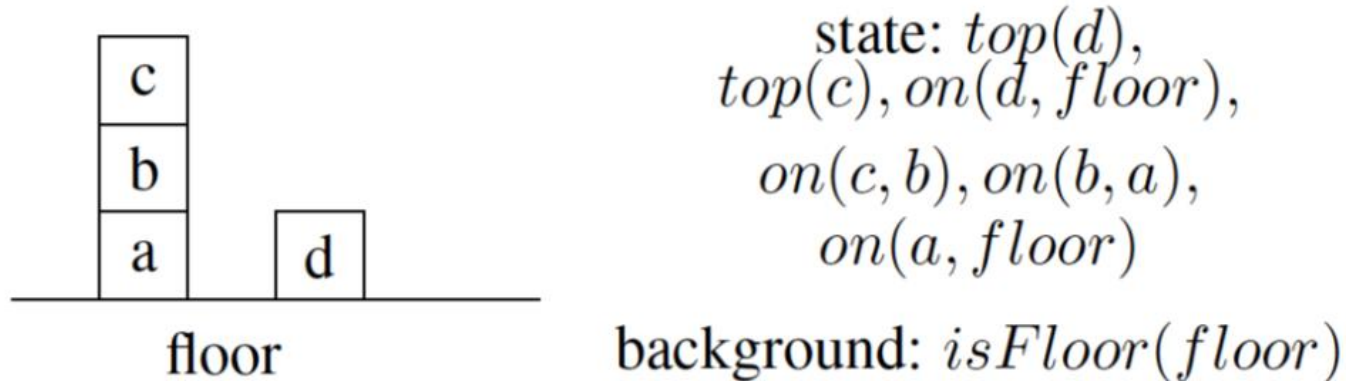


Figure: Block Manipulation

# Problem Statement

A Markov Game is defined by a tuple  $(\mathcal{N}, \mathcal{S}, \{\mathcal{A}^i\}_{i \in \mathcal{N}}, \mathcal{P}, \{R^i\}_{i \in \mathcal{N}}, \gamma)$ , where  $\mathcal{N}$  denotes the set of  $N$  agents.  $\mathcal{A}^i$  denotes the action space of agent  $i$ .  $\mathcal{P}$  is transition probability.  $\mathcal{S}$  denotes the state space.

Social Rules is defined by a map,  $F : S \times \{\mathcal{A}\} \rightarrow \{(0, 1)\}^n$ . where  $\{\mathcal{A}\}$  is action space  $(a_1, a_2, \dots, a_n)$ .

Map  $F$  can be represented logic formula, e.g. First-order logic.

# Problem Statement

Under the limit of social rules, the joint policy  $\pi$  will be  $\pi(a|s) := \prod_{i \in \mathcal{N}} \pi^i(a^i|s) |_{F(s)}$ .  
Optimize Value Function and get optimal  $\pi^*$ :

$$R = \sum_t \gamma^t R_k^t, R_k = \{0, -1, 1\}$$

$$V_{\pi^i, \pi^{-i}}^i(s) := \mathbb{E} \left[ \sum_{t \geq 0} \gamma^t R^i(s_t, a_t, s_{t+1}) \middle| a_t^i \sim \pi^i(\cdot | s_t), s_0 = s \right],$$

where  $i$  represents the indices of all agents in  $N$  except agent  $i$ .

# Experiment

- Tic-Tac-Toe

- State  $S: (i, j) \rightarrow \{-1, 0, 1\}$

- Mine(i,j)  $\rightarrow 1$
- Opponent(i,j)  $\rightarrow -1$
- Empty(i,j)  $\rightarrow 0$


- Action A:

Place(x,y)

- Reward R:

$$R = \sum_t \gamma^t R_k^t, R_k = \{0, -1, 1\}$$

- Auxiliary predicate: Invented(i) ,  $i = 1, 2, 3, 4$



# Experiment ---Tic-Tac-Toe

- Invalid action problem:
  - Place(x,y) :- Mine(x,y), Mine(x,z)
- One solution:
  - Give penalty for invalid action

$$R = \sum_t \gamma^t p_t R_k^t, R_k = \{0, -1, 1\}$$


# Tic-Tac-Toe: Results

- State

```
([[ 1.,  0., -1.],  
 [ 0.,  0.,  0.],  
 [-1.,  0.,  1.]])
```

- Logic rules:

```
place(X,Y):-invented4(X),invented4(Y)  
invented4(X):-empty(X,X),succ(X,Y)
```




# Tic-Tac-Toe: Results

- State:

```
[[-1.,  0.,  0.],  
 [ 0.,  0.,  0.],  
 [ 1.,  0., -1.]]
```

- Logic rules:

```
place(X,Y):-empty(X,X),empty(Y,Y)
```






# Tic-Tac-Toe: Results

- State:

```
[[ 0.  0.  0.]  
 [ 0.  1.  0.]  
 [-1.  0. -1.]]
```

- Logic rules:

```
place(X,Y):-empty(Z,Y),succ(Y,X)
```



# Future Work

- Using more powerful logic representation: SDD
- Two more experiments: Escalator Traffic Rules
- Combine social rules learning and multi-agent learning

SDD: Sentential Decision Diagram

